

Hallucination Detection and Evaluation in Multimodal Large Language Models

Xiangrui Yue

May 24, 2026

Abstract

Multimodal Large Language Models (MLLMs) have demonstrated strong capabilities in visual question answering and image-text reasoning, yet they remain susceptible to hallucinations: outputs that are inconsistent with the visual input, contradict established world knowledge, or contain flaws in the reasoning process. This report presents a framework for the automatic detection and evaluation of hallucinations in MLLMs. I adopt a three-category taxonomy covering faithfulness, factuality, and logical hallucinations, and evaluate four representative models (GPT-5.4-mini, Gemini 2.5 Flash, Qwen3.5-35B-A3B, Qwen3-VL-235B-A22B) across three benchmarks (POPE, MathVista, VQA-RAD) using two detection methods (a rule-based pipeline and an MLLM-Judge protocol), yielding two findings. (1) Language priors structurally override visual perception. POPE’s three splits expose the confounding effects of two modality priors: the parametric bias of the visual encoder (P_V) and the statistical co-occurrence regularities acquired by the LLM from text-only pre-training (P_L), which influence the output through independent attention pathways. When visual signals are ambiguous, P_L suppresses residual visual evidence (Popular HR 2.40%–6.20%); when P_V and P_L fail jointly, conceptual boundary blur and semantic conflation combine to produce the most severe false positives (Adversarial HR 4.80%–8.90%). This mechanism also explains why faithfulness hallucinations account for over 96% of errors in VQA-RAD: fluent medical knowledge cannot compensate for inadequate radiological visual grounding. (2) The effect of Chain-of-Thought prompting depends on the model’s baseline: CoT substantially reduces hallucinations on weaker baselines (GPT-5.4-mini –10.7 pp, Gemini 2.5 Flash –6.7 pp) but increases them on stronger ones (Qwen3-VL +7.7 pp). I attribute this two-sided effect to four interacting mechanisms: structured decomposition and visual grounding help weaker models correct reasoning and perception errors, whereas exposure effects and commitment errors make stronger models explicitly articulate latent visual misreadings and continue reasoning from erroneous premises. Five independent error analyses (human alignment with Cohen’s $\kappa = 0.70$, cross-judge consistency between GPT-5.5 and Claude Opus 4.7 with $\kappa = 0.78$, threshold sensitivity, length-bias assessment, and cross-model overlap) confirm the robustness of the evaluation framework; code and data are publicly available at <https://github.com/lexiangrui/Hallucination-of-MLLM>.

Contents

1	Introduction	3
1.1	Task Description	3
1.2	Evaluated Models	3
1.3	Evaluation Datasets	4
2	Hallucination Definition and Taxonomy	4
2.1	Faithfulness Hallucination (Visual Inconsistency)	5
2.2	Factuality Hallucination (Factual Inconsistency)	6
2.3	Logical Hallucination (Reasoning Error)	6

3	Experimental Design	7
3.1	Detection Methods and Evaluation Metrics	7
3.2	Experimental Conditions	9
4	Experiment 1: POPE Object Probing	9
4.1	Results	9
4.2	Analysis	9
4.3	Failure Case Analysis	12
5	Experiment 2: MathVista Mathematical Reasoning	14
5.1	Results	14
5.2	Analysis	14
5.3	Failure Case Analysis	16
6	Extension: VQA-RAD Medical VQA	18
6.1	Results	18
6.2	Result Analysis	18
6.3	Failure Case Analysis	19
7	Error Analysis	21
7.1	Human Alignment Verification (Human-as-Judge)	21
7.2	Judge Model Consistency Check	23
7.3	Threshold Sensitivity Analysis	24
7.4	Length Bias Analysis	24
7.5	Cross-Model Hallucination Consistency	25
8	Limitations and Future Work	26
8.1	Limitations	26
8.2	Future Work	26
	References	28
A	Dataset Details and Prompt Templates	31
A.1	Dataset Details	31
A.2	Response Generation Prompts	32
A.3	MLLM Judge Prompts	33

1 Introduction

1.1 Task Description

Multimodal Large Language Models (MLLMs) have demonstrated strong capabilities in visual question answering and image-text reasoning, yet they remain susceptible to hallucinations: outputs that are inconsistent with the visual input, contradict established world knowledge, or contain flawed reasoning. This report investigates hallucination detection and evaluation in MLLMs. The assessment requirements are summarized as follows:

1. **Hallucination definition and classification:** define at least two types of hallucinations, provide clear explanations, and supply no fewer than two examples for each type.
2. **Automated detection methods and evaluation metrics:** construct at least two detection methods (e.g., a rule-based method and an MLLM-Judge protocol) and design no fewer than two evaluation metrics.
3. **Experimental evaluation (main task):** evaluate at least two models, including one closed-source and one open-source, on general visual question answering (e.g., POPE) and one of the recommended image-text reasoning benchmarks (MathVista or ChartQA).
4. **Human alignment verification (Human-as-Judge):** construct a small-scale human-annotated set of 20–50 samples as the gold label, compare automated detection results with human judgments, and analyze inconsistent cases.
5. **Analysis requirements:** include an analysis of CoT’s impact on hallucinations, a cross-model comparison, and in-depth analysis of no fewer than two typical failure cases.
6. **Bonus:** extend the evaluation to other image-text tasks or specific domains (e.g., OCR, medical, or remote sensing), or conduct a cross-task generalization analysis.

1.2 Evaluated Models

I evaluate four representative MLLMs that cover both closed-source and open-source models and span a range of architectures (dense and Mixture-of-Experts) and parameter scales (from 3 B to 235 B activated). Two additional models serve as MLLM Judges in the detection pipeline and in the cross-judge consistency check (Appendix A). Table 1.1 summarizes the key configurations.

Table 1.1: Model configurations. Four models are evaluated; two additional models serve as MLLM Judges. Architecture lists the total parameter count for MoE models; the active counts are 3 B (Qwen3.5-35B-A3B) and 22 B (Qwen3-VL-235B-A22B).

Model	Role	Architecture	Access Type	API SDK
GPT-5.4-mini (OpenAI, 2026b)	Evaluated	Dense	Closed-source	OpenAI
Gemini 2.5 Flash (Google DeepMind, 2026)	Evaluated	Dense	Closed-source	OpenAI
Qwen3.5-35B-A3B (Qwen Team, 2025a)	Evaluated	MoE 35 B	Open-source	OpenAI
Qwen3-VL-235B-A22B (Qwen Team, 2025b)	Evaluated	MoE 235 B	Open-source	OpenAI
GPT-5.5 (OpenAI, 2026a)	Judge	Dense	Closed-source	OpenAI
Claude Opus 4.7 (Anthropic, 2026)	Judge	Dense	Closed-source	Anthropic

1.3 Evaluation Datasets

I adopt three standard benchmarks that respectively target object-level perception, mathematical visual reasoning, and medical radiology VQA. All datasets are accessed from Hugging Face.¹ Table 1.2 provides a summary; detailed descriptions are deferred to Appendix A.

Table 1.2: Dataset overview. POPE samples are subsampled per split for this experiment; MathVista uses the full TESTMINI split and VQA-RAD uses the official test split.

Dataset	Task Type	Samples
POPE (Li et al., 2023)	Object existence probing	3,000 (3 splits $\times$ 1,000)
MathVista (Lu et al., 2024)	Mathematical visual reasoning	1,000 (testmini)
VQA-RAD (Lau et al., 2018)	Medical radiology VQA	451 (test split)

2 Hallucination Definition and Taxonomy

Hallucination remains a central challenge for MLLMs. Although different studies adopt varying terminology and perspectives, a broad consensus has emerged on the underlying concepts. This section reviews three representative surveys (Bai et al., 2025b; Chen et al., 2026; Liu et al., 2024a) and then presents the taxonomy adopted in this work.

Bai et al. (2025b) define hallucination in MLLMs as generated text that is inconsistent with the corresponding visual content, emphasizing cross-modal inconsistency. They also trace the concept back to LLMs, distinguishing factual hallucinations (deviations from real-world facts) from faithfulness hallucinations (deviations from the input context or self-consistency). Their object-centric taxonomy comprises three levels: category (fabricating or misidentifying objects), attribute (correct category but incorrect properties such as color, shape, or count), and relation (correct objects and attributes but incorrect spatial or interactive relations).

Chen et al. (2026) propose a broader unified definition, treating hallucination as any inconsistency between the generated content and either the input conditions or established world knowledge. Their taxonomy extends the faithfulness-factuality dichotomy: faithfulness hallucinations are stratified by visual granularity into object-level, attribute-level, and scene-level errors, forming a progression from shallow perception to deeper scene understanding; factuality hallucinations are divided into commonsense errors and domain-specific errors.

Liu et al. (2024a) restrict hallucination to the visual domain, defining it as inconsistency between the factual visual content of the image and the generated text. They adopt a dual-perspective taxonomy: a cognitive perspective distinguishing judgment hallucinations (yes/no contradictions with visual facts) from descriptive hallucinations (unfaithful free-form generation), and a semantic perspective distinguishing object, attribute, and relation hallucinations.

Synthesizing these frameworks and given that the experiments in this report primarily target visual question answering and image-text reasoning, I adopt the following definition and three-category taxonomy.

An MLLM hallucination is a model response that is inconsistent with the input visual content, contradicts verifiable real-world knowledge, or contains flaws in the reasoning process.

Table 2.1 situates this definition relative to the three surveys above. The three categories are detailed below.

¹POPE: <https://huggingface.co/datasets/lmms-lab/POPE>; MathVista: <https://huggingface.co/datasets/AI4Math/MathVista>; VQA-RAD: <https://huggingface.co/datasets/flaviagiammarino/vqa-rad>.

Table 2.1: Comparison of MLLM hallucination taxonomies across three representative surveys and the framework adopted in this work. The bottom row (shaded) marks the taxonomy used throughout this report.

Source	Scope	Primary axis	Sub-categories
Bai et al. (2025b)	Cross-modal inconsistency	Factual vs. faithfulness	Category / attribute / relation
Chen et al. (2026)	Generalized inconsistency	Faithfulness vs. factuality	Faithfulness: object / attribute / scene; Factuality: commonsense / domain
Liu et al. (2024a)	Visual domain only	Judgment vs. descriptive	Object / attribute / relation
This work	Visual + knowledge + reasoning	Three categories	Faithfulness / factuality / logical


2.1 Faithfulness Hallucination (Visual Inconsistency)

Faithfulness hallucination refers to a model response that is inconsistent with the visual content of the input image. Depending on the level of visual granularity, this category is further divided into three sublevels.

At the object level, the model claims the presence of objects that do not exist in the image, omits objects that are actually present, or misidentifies object categories. The POPE benchmark systematically detects this type of hallucination through binary probing questions (Figure 2.1a).

At the attribute level, the model correctly identifies the object category but incorrectly describes its properties such as color, shape, material, or pose (Figure 2.1b).

At the scene level, spatial relations, interactive events, or other scene-level descriptions are inconsistent with the image, even though all individual objects are correctly identified (Figure 2.1c).



An image of a husky lying on a couch. Object-level faithfulness hallucinations include category substitution (reporting a cat instead of a dog), object omission (claiming no animal is present), and existence misjudgment (affirming the presence of a non-existent object).

(a) Object level.



A round clock mounted on a pole beside a city street. The model correctly identifies the clock but may hallucinate incorrect attributes, such as describing it as square or misreporting the color and hand positions.

(b) Attribute level.



A kitchen countertop with a knife on a cutting board, a tequila bottle to the left of a blender, and a glass to the right. Scene-level hallucinations involve incorrect spatial relations or fabricated interactions between correctly identified objects.

(c) Scene level.

Figure 2.1: Faithfulness hallucination examples at three levels of visual granularity.


The central detection question for this category is: is the output consistent with the image?

2.2 Factuality Hallucination (Factual Inconsistency)

Factuality hallucination refers to a model response that contradicts established, verifiable real-world knowledge.

At the commonsense level, the error involves everyday knowledge, where the model’s visual perception is roughly correct but the output contradicts well-known facts (Figure 2.2a).

At the domain-specific level, the error involves specialized professional knowledge. Such hallucinations are particularly common in the VQA-RAD dataset (Figure 2.2b).



An image showing zebras and giraffes. Commonsense factuality hallucinations arise when the model’s roughly correct visual perception leads to outputs that contradict basic biological knowledge, such as misidentifying a zebra as a painted horse.

(a) Commonsense level.



A chest X-ray with no obvious abnormalities in the lung fields. Domain-specific factuality hallucinations occur when the model fabricates findings such as pleural effusion or cardiomegaly that contradict the actual radiological evidence.

(b) Domain-specific level.

Figure 2.2: Factuality hallucination examples at the commonsense and domain-specific levels.


The central detection question for this category is: is the output consistent with world knowledge?

2.3 Logical Hallucination (Reasoning Error)

Logical hallucination refers to a model response with flaws in the reasoning process, where the conclusion does not logically follow from the visual evidence or from its own established premises.


In a deductive reasoning error, the model correctly reads the visual information but reaches a contradictory conclusion through faulty inference. Such errors, particularly numerical reasoning failures, are the primary source of logical hallucinations in MathVista (Figure 2.3a).

In a causal reasoning error, the model infers causal relationships that contradict the visual evidence (Figure 2.3b).



A bar chart showing government expenditure percentages for Czechia, Paraguay, and Laos. The model correctly reads the bar values but errs in the comparison step, producing a conclusion that contradicts its own extracted evidence.

(a) Deductive reasoning.



A food web showing predation relationships among field mice, snakes, hawks, grasshoppers, frogs, and grass. Causal reasoning hallucinations arise when the model incorrectly traces causal chains, reversing the expected direction of population effects.

(b) Causal reasoning.

Figure 2.3: Logical hallucination examples in deductive and causal reasoning.

The central detection question for this category is: is the reasoning valid?

3 Experimental Design

3.1 Detection Methods and Evaluation Metrics

3.1.1 POPE: Object-Existence Probing

Object hallucination in POPE. Li et al. (2023) define object hallucination as the phenomenon in which a vision-language model generates descriptions or answers containing objects that are inconsistent with or absent from the target image. POPE operationalizes this notion as a binary probing task: for each image, the benchmark issues yes/no questions of the form “Is there an X in the image?”, and the model commits an object hallucination whenever it answers “yes” while the ground-truth label is “no.” The event is a false positive in the binary-classification view and instantiates the object-level faithfulness category of the taxonomy.

The metrics treat “yes” (object present) as the positive class, so object hallucinations correspond to false positives. The core metrics are summarized in Table 3.1.

Table 3.1: POPE evaluation metrics. Hallucination Rate is the core hallucination metric.

Metric	Formula	Description
Accuracy	$(TP + TN)/N$	Overall correctness
Precision	$TP/(TP + FP)$	Accuracy when predicting Yes
Recall	$TP/(TP + FN)$	Detection rate for actual Yes
F1	$2PR/(P + R)$	Harmonic mean of Precision and Recall
Yes Ratio	N_{yes}/N	Response bias indicator
Hallucination Rate	FP/N	Core metric; equals object hallucination rate

Prompt and post-processing. The detection method for POPE is rule-based matching. During response generation, the model receives a prompt containing only the question and an instruction to answer Yes or No, ensuring concise and parseable responses (the complete prompt template is provided in Appendix A). The answer post-processing algorithm follows the official evaluation script of the RUCAIBox/POPE repository;² the pseudocode is given in Algorithm 1.

Algorithm 1 POPE answer normalization

Require: Model response r , a free-form text string

Ensure: Normalized answer $a \in \{\text{“yes”}, \text{“no”}\}$

- 1: $s \leftarrow$ first sentence of r *// split by “.”*
 - 2: remove all commas from s
 - 3: $W \leftarrow$ split s by whitespace into a set of words
 - 4: **if** “No” $\in W$ **or** “not” $\in W$ **or** “no” $\in W$ **then**
 - 5: **return** “no”
 - 6: **end if**
 - 7: **return** “yes”
-

3.1.2 MathVista: Mathematical Visual Reasoning

Detection method. The detection method for MathVista employs an MLLM Judge following the MMHal-Bench evaluation protocol (Sun et al., 2024), which adopts a 0–6 Likert scale that jointly assesses the informativeness and the hallucination level of each response. A threshold of 3 separates the two regimes: scores below 3 indicate hallucination, and scores at or above 3 indicate no hallucination. GPT-5.5 serves as the Judge in this implementation. The judge’s

²<https://github.com/RUCAIBox/POPE/blob/main/evaluate.py>

system prompt defines the scoring rubric, the hallucination criteria, and a three-way type classification, and is supplemented by five few-shot examples for consistency. Whenever the judge assigns a score below 3, it additionally labels the hallucination type (faithfulness, factuality, or logical). The complete prompt templates for both response generation and judge scoring are provided in Appendix A.




Figure 3.1: MLLM Judge evaluation framework for MathVista: the judge receives a question, a model response, and a scoring rubric, and produces a score (0–6) with an optional hallucination type label.

Prompt construction. During response generation, prompts are dynamically constructed from sample metadata to produce format-specific instructions. In Direct mode, the prompt consists of the question text concatenated with format instructions. In CoT mode, the prompt further requires step-by-step reasoning with the final answer marked by “Final answer:”. Both modes are evaluated to analyze the effect of explicit reasoning on hallucination behavior.

The evaluation metrics are shown in Table 3.2.

Table 3.2: MathVista evaluation metrics. Hallucination Rate is the core hallucination metric.

Metric	Definition	Description
Hallucination Rate	$N_{\text{score}<3}/N_{\text{total}}$	Lower is better
Average Score	$\sum \text{score}_i / N$	Overall response quality
Type-specific count	By faithfulness/factuality/logical	Hallucination source breakdown

3.1.3 VQA-RAD: Medical Radiology VQA

Medical adaptation. The detection method for VQA-RAD reuses the MLLM Judge protocol (0–6 Likert scale, threshold 3) with adaptations for the medical setting (see Appendix A). The judge’s task description is replaced with a radiology-specific scenario description, requiring

the judge to verify consistency with both the image’s visual content and established medical and anatomical knowledge. During response generation, a radiologist expert role prompt is adopted following the role-setting strategy of Wu et al. (2024). The three-way hallucination type classification (faithfulness, factuality, logical) is always enabled, and hallucination rates are additionally stratified by closed-ended and open-ended answer types.

The full 451-sample test split is used, comprising 251 closed-ended (yes/no) and 200 open-ended (free-text) questions. The evaluation metrics include Hallucination Rate, Average Score, type-specific hallucination counts, and hallucination rates stratified by answer type.

3.2 Experimental Conditions

All four models are accessed through APIs with a unified random seed of 42. Temperature is kept at each platform’s default, and the maximum token length is uncapped. POPE samples 1,000 items per split, MathVista uses the full testmini set (1,000 items), and VQA-RAD uses the test split (451 items). GPT-5.5 serves as the MLLM Judge throughout, and a human alignment study is conducted as part of the error analysis.

Table 3.3 summarizes the correspondence among detection methods, datasets, and metrics.

Table 3.3: Overview of method–dataset–metric correspondence.

Dataset	Detection Method	Annotation Scope	Core Metrics
POPE	Rule-based	Faithfulness	HR, F1
MathVista	MLLM Judge	All three types	HR, Avg Score
VQA-RAD	MLLM Judge	All three types	HR, Avg Score, Closed/Open HR
Human	Manual annotation	All three types	Kappa, F1, Type Agreement

4 Experiment 1: POPE Object Probing

4.1 Results

Table 4.1 reports all evaluation metrics for the four models across the three POPE splits. Columns are grouped by split, each containing Accuracy, Precision, Recall, F1, Yes Ratio, Hallucination Rate (HR), and the four confusion matrix entries (TP/TN/FP/FN).

4.2 Analysis

4.2.1 Cross-Model Comparison

The four models exhibit distinct decision styles that can be characterized along two dimensions: hallucination control (HR) and response bias (Yes Ratio). Figure 4.1 visualizes the HR of all four models across the three splits.

Qwen3-VL-235B achieves the lowest HR on every split (0.30%, 2.40%, 4.80%) with a stable Yes Ratio (≤ 0.47), reflecting a strong visual encoder combined with a balanced decision style.

Gemini 2.5 Flash shows balanced Precision–Recall and is the only model whose HR *decreases* from Popular (5.70%) to Adversarial (5.10%), indicating greater robustness to near-synonym distractors than to co-occurrence distractors.

Qwen3.5-35B-A3B attains the highest Recall on every split (≥ 0.87) but at the cost of the panel’s highest Yes Ratio (peaking at 0.529 on Adversarial) and worst Adversarial HR (8.90%), exemplifying a permissive decision style.

GPT-5.4-mini adopts the opposite, conservative strategy: the lowest Yes Ratio (0.414 on Random) suppresses HR on easier splits but produces the steepest cross-split degradation

Table 4.1: Comprehensive POPE results across three splits. Best values per metric are highlighted in pink.

Split	Metric	GPT-5.4-mini	Gemini 2.5 Flash	Qwen3.5-35B	Qwen3-VL-235B
Random	Accuracy	0.9010	0.9160	0.9260	0.9200
	Precision	0.9831	0.9814	0.9754	0.9929
	Recall	0.8156	0.8477	0.8737	0.8457
	F1	0.8916	0.9097	0.9218	0.9134
	Yes Ratio	0.414	0.431	0.447	0.425
	HR	0.70%	0.80%	1.10%	0.30%
	TP/TN/FP/FN	407/494/7/92	423/493/8/76	436/490/11/63	422/498/3/77
Popular	Accuracy	0.8570	0.8620	0.8730	0.8970
	Precision	0.8836	0.8800	0.8750	0.9459
	Recall	0.8216	0.8377	0.8697	0.8417
	F1	0.8515	0.8583	0.8724	0.8908
	Yes Ratio	0.464	0.475	0.496	0.444
	HR	5.40%	5.70%	6.20%	2.40%
	TP/TN/FP/FN	410/447/54/89	418/444/57/81	434/439/62/65	420/477/24/79
Adversarial	Accuracy	0.8280	0.8730	0.8520	0.8750
	Precision	0.8385	0.8924	0.8318	0.8979
	Recall	0.8116	0.8477	0.8818	0.8457
	F1	0.8248	0.8695	0.8560	0.8710
	Yes Ratio	0.483	0.474	0.529	0.470
	HR	7.80%	5.10%	8.90%	4.80%
	TP/TN/FP/FN	405/423/78/94	423/450/51/76	440/412/89/59	422/453/48/77




Figure 4.1: Hallucination rates of four MLLMs across the three POPE splits (Random, Popular, Adversarial).

(Adversarial HR = 7.80%, F1 = 0.8248, both panel-worst), as the “default to No” shortcut fails under adversarial pressure.

4.2.2 Three Splits Detect Three Error Mechanisms

Monotonic difficulty gradient. From Random to Popular to Adversarial, the Hallucination Rate of all four models increases monotonically (Table 4.1 and figure 4.1), confirming the validity of POPE’s three-tier negative-sampling design. Beyond a difficulty ranking, however, the three splits also expose three qualitatively distinct error mechanisms that map cleanly onto the faithfulness taxonomy.

Random split: object-level perception. Negative objects are sampled uniformly, so most candidates are visually unrelated to the scene (e.g., asking about a car in a kitchen photo). All four models achieve HR between 0.30% and 1.10%, indicating that the baseline ability to reject obviously absent objects is near saturation. This split therefore measures raw object-level visual perception, with negligible interference from language priors.

Popular split: resistance to linguistic priors. Negative objects are high-frequency co-occurring categories that are statistically associated with the scene but absent from the specific image (e.g., asking about a backpack in an outdoor eating scene). HR rises sharply to 2.40%–6.20%, roughly 5–8 times the Random level. The surge isolates the contribution of language-prior bias: when the visual evidence is ambiguous, models fall back on the statistical co-occurrence patterns learned during pretraining. Case 2 in Section 4.3.2 provides a direct instance of this mechanism.

Adversarial split: near-synonym semantic discrimination. Negative objects are semantically similar to objects that do appear in the image (e.g., asking about a bear when only a teddy bear is present). HR rises further to 4.80%–8.90% and the model can no longer rely on scene-level associations: the visual encoder must resolve fine-grained semantic boundaries. Case 1 in Section 4.3.1 shows that this is the hardest of the three failure modes, because the underlying cause, confusion of conceptual hierarchies, is not addressable by scale alone.

4.2.3 Root Cause: The Confounding Effects of Dual Modality Priors

The architectural asymmetry of MLLMs. Current mainstream MLLMs all adopt the modular architecture of “vision encoder + projector + LLM backbone”: the vision encoder (typically a CLIP or SigLIP ViT) converts an image into feature vectors, which an MLP projector then maps into the LLM’s word embedding space; the resulting visual tokens are concatenated with text tokens and fed into the LLM for autoregressive generation. This architecture harbors a fundamental asymmetry: the parameter count and pretraining data scale of the LLM backbone far exceed those of the visual pathway. Taking Qwen3-VL-235B-A22B as an example, its vision encoder is SigLIP2-SO-400M (400M parameters) while the language backbone contains 235B total parameters (22B active)—a difference of roughly two orders of magnitude ($588\times$ in total parameters) (Bai et al., 2025a). In the shared self-attention computation where visual and text tokens interact, this scale disparity grants the language modality a structural advantage at the parameter level.

Two independent confounding factors. Zhou et al. (2025) demonstrated through a structural causal model in CausalMM (ICLR 2025) that this architectural asymmetry introduces two independent confounding factors (Figure 4.2): (1) the **visual modality prior** P_V , originating from the parametric bias of the vision encoder—the training distribution of CLIP/SigLIP determines the encoder’s sensitivity or blindness to specific visual patterns; and (2) the **linguistic modality prior** P_L , originating from the statistical co-occurrence regularities that the LLM has acquired from billions of tokens of text-only pretraining. These two priors influence the final output through distinct attention pathways: P_V distorts the attention distribution within the vision encoder, thereby contaminating the visual embeddings passed to the LLM;

P_L activates pathways in the LLM’s multi-head self-attention that are consistent with linguistic statistics, and its effect is most pronounced at the shallow-to-middle layers where visual tokens first interact with text tokens.




Figure 4.2: Structural Causal Model (SCM) of MLLM inference proposed by CausalMM. Blue nodes (V , L) represent attention mechanisms; red nodes (P_V , P_L) are modality priors acting as confounders; gray nodes are image (I), visual embeddings (T_V), text embeddings (T_L), and output (O). Red paths indicate the confounding effects of P_V and P_L on attention and output. Source: Zhou et al. (2025), ICLR 2025.

A unified causal account of the three POPE error mechanisms. Returning to the experimental results (Section 4.2.2), the three error patterns correspond precisely to how these two priors interact under different conditions. In the Random split, negative objects are unrelated to the scene and P_L is barely engaged; the residual hallucination rate of 0.30%–1.10% stems from the baseline bias of P_V alone, reflecting the upper bound of the vision encoder’s perceptual capacity. The Popular split represents the P_L effect in its purest form—when visual evidence is ambiguous, the linguistic co-occurrence prior (“outdoor + eating $\rightarrow$ backpack”) suppresses weak visual signals through the LLM’s shallow-to-middle layer self-attention, causing the hallucination rate to surge to 2.40%–6.20%. The Adversarial split exhibits a joint failure of P_V and P_L : the vision encoder (P_V bias) cannot reliably separate “teddy bear” and “bear” at the fine-grained semantic boundary in feature space, and the LLM (P_L) subsequently merges the two at the semantic level, triggering a false positive. This also explains why increasing model scale cannot resolve Adversarial hallucinations (Qwen3-VL-235B still exhibits 4.80%)—scaling primarily strengthens P_L ’s statistical precision, without altering the intrinsic fuzziness of conceptual boundaries within P_V ’s feature space.

4.3 Failure Case Analysis

4.3.1 Case 1: Near-synonym Semantic Confusion (Adversarial)

This case instantiates the near-synonym semantic confusion mechanism identified in Section 4.2.2. Table 4.2 and figure 4.3 present a sample from the Adversarial split. The image contains a teddy bear, a person, and a book. The question asks whether there is a bear in the image, and the ground truth is No. All four models answer Yes, indicating a systematic error rather than an individual model deficiency.

The root cause is near-synonym semantic confusion. The image contains a teddy bear, an independent category in the COCO label set (Lin et al., 2014), while the question asks about a bear (a real animal). All four models associate “teddy bear” with “bear” at the semantic level, triggering a positive judgment upon seeing the toy. This is essentially conceptual hierarchy confusion: models conflate visually similar objects with semantically identical ones. This error cannot be resolved by increasing model scale or improving visual encoder resolution, as it concerns the clarity of conceptual boundaries in semantic representations rather than the discriminability of low-level visual features.

Table 4.2: POPE Case 1: sample information and model responses.

Sample Information	
Image	COCO_val2014_000000567886
Objects present	teddy bear, person, book
Objects absent	bear, car, tv
Question	Is there a bear in the image?
Ground Truth	No
Model Responses	
GPT-5.4-mini	Yes
Gemini 2.5 Flash	Yes
Qwen3.5-35B-A3B	Yes
Qwen3-VL-235B	Yes



Figure 4.3: POPE Case 1: a teddy bear on a couch.

4.3.2 Case 2: Linguistic Prior Bias (Popular)

This case exemplifies the linguistic-prior bias mechanism isolated by the Popular split (Section 4.2.2). Table 4.3 and figure 4.4 present a sample from the Popular split. The image shows a person eating a sandwich and a hot dog outdoors. The question asks whether there is a backpack in the image, and the ground truth is No. Again, all four models answer Yes.

Table 4.3: POPE Case 2: sample information and model responses.

Sample Information	
Image	COCO_val2014_000000569839
Objects present	person, sandwich, hot dog
Objects absent	car, bottle, backpack
Question	Is there a backpack in the image?
Ground Truth	No
Model Responses	
GPT-5.4-mini	Yes
Gemini 2.5 Flash	Yes
Qwen3.5-35B-A3B	Yes
Qwen3-VL-235B	Yes



Figure 4.4: POPE Case 2: an outdoor dining scene.

The root cause is linguistic prior bias. In outdoor scenes with people eating, “backpack” is a high-frequency co-occurring object. The statistical prior in the language module (outdoor + person + eating $\rightarrow$ high backpack probability) overrides the visual module’s actual perception, causing the model to produce an affirmative answer. This is precisely the mechanism that the Popular split exploits: models fill visual uncertainty with linguistic priors from training data. The contrast between the two cases further confirms that the Popular and Adversarial splits target distinct error mechanisms, with Popular primarily exploiting linguistic priors and Adversarial primarily testing semantic discrimination.

5 Experiment 2: MathVista Mathematical Reasoning

5.1 Results

Table 5.1 summarizes the core MathVista evaluation metrics for all four models under Direct and CoT modes.

Table 5.1: MathVista Direct and CoT combined results. Δ values are CoT minus Direct; negative values indicate fewer hallucinations under CoT, and positive net improvement indicates better overall performance. Best values are highlighted in pink.

Model	Direct HR	CoT HR	Δ HR	Direct Avg	CoT Avg	Δ Avg	Δ Faith.	Δ Fact.	Δ Logical	Net improvement
GPT-5.4-mini	45.1%	34.4%	−10.7	3.31	4.35	+1.04	+8	−10	−105	+107
Gemini 2.5 Flash	34.3%	27.6%	−6.7	4.46	4.81	+0.35	−66	−8	+7	+67
Qwen3.5-35B-A3B	13.3%	15.5%	+2.2	5.00	5.34	+0.34	+26	−3	−1	−22
Qwen3-VL-235B	16.3%	24.0%	+7.7	4.85	4.94	+0.09	+41	0	+36	−77




Figure 5.1: Hallucination type composition among hallucinated samples under MathVista Direct and CoT modes.

5.2 Analysis

5.2.1 Differences in Hallucination Type Distribution

The hallucination-type composition in Figure 5.1 shows that faithfulness hallucinations dominate most configurations. Except for GPT-5.4-mini in Direct mode, faithfulness accounts for more than 57% of hallucinated samples in the remaining seven model-mode configurations, peaking at 81% for Gemini Direct. This indicates that MathVista’s core bottleneck is that the model must first “see” the numbers, labels, and spatial relations in the chart clearly; reasoning comes only after that.

The main exception is GPT-5.4-mini in Direct mode, where logical hallucinations account for 55.0% of hallucinated samples. This does not reflect unusually poor reasoning ability alone. Rather, it reflects its extremely terse response style: the median Direct response length is only one word, so the judge receives almost no reasoning process to evaluate. With no reasoning chain available, wrong final answers are naturally categorized as logical failures. Once CoT forces the model to unfold its derivation, logical hallucinations drop sharply by 105 cases and the type distribution returns to a level closer to the other models (41.6% logical). Factuality hallucinations remain below 4% in all models and modes, showing that MathVista errors rarely come from missing world knowledge; they concentrate instead in visual perception and reasoning.




Figure 5.2: Comparison of hallucination rates between Direct and CoT modes. Δ values indicate the change.

5.2.2 Comparison Between Direct and CoT Modes

CoT effects on MathVista divide cleanly according to Direct-mode baseline capability. Models with high Direct HR (GPT-5.4-mini at 45.1%, Gemini at 34.3%) improve substantially after CoT is introduced, whereas models with low Direct HR (Qwen3.5 at 13.3%, Qwen3-VL at 16.3%) degrade, as shown in Figure 5.2.

CoT significantly improves weaker-baseline models. GPT-5.4-mini is the largest beneficiary: HR drops by 10.7 percentage points and average score increases by 1.04, both the largest improvements among all models. This improvement comes primarily from the sharp reduction in logical hallucinations (from 248 to 143). In Direct mode this model outputs only option letters, exposing incorrect answers directly to the judge. CoT forces the model to organize its mathematical derivation step by step, thereby correcting many errors that a reasoning chain can catch. Gemini 2.5 Flash improves through a different pathway: its gain mainly comes from fewer faithfulness hallucinations (from 278 to 212). CoT requires the model to explicitly cite visual evidence while reasoning, which forces a more careful examination of chart content and reduces visual misreadings.

CoT introduces additional risk for stronger-baseline models. Qwen3.5-35B-A3B shows a 2.2 percentage-point increase in HR. Because this model already possesses strong mathematical reasoning in Direct mode, CoT provides little marginal benefit. Instead, by forcing explicit visual description, it exposes blind spots in visual perception, causing faithfulness hallucinations to rise from 98 to 124.

Qwen3-VL-235B shows a more pronounced pattern: HR jumps from 16.3% to 24.0%, with a net degradation of 77 samples. Of these, 116 Direct-mode correct samples produce hallucinations under CoT. Faithfulness hallucinations rise from 113 to 154, indicating that forced visual articulation systematically exposes perceptual weaknesses.

5.2.3 Underlying Mechanisms of the CoT Effect

This phenomenon aligns with recent findings. Liu et al. (2025) show through attention analysis that as reasoning chains grow longer, the model’s attention to visual input weakens (attention drift), gradually shifting from image content to linguistic priors, leading to “more thinking,

less seeing.” Kancheti et al. (2025) report, after evaluating 17 models, that CoT consistently degrades visual spatial reasoning, with models “hallucinating visual details from textual priors.” The substantial increase in faithfulness hallucinations for both Qwen models is consistent with both observations.

Taken together, the results suggest that CoT both improves and introduces hallucinations through four simultaneous mechanisms. The first two explain improvements on weaker baselines, whereas the last two explain degradation on stronger baselines.

Structured decomposition. Wei et al. (2022) identify the central benefit of CoT: decomposing a complex multi-step problem into controllable intermediate steps. GPT-5.4-mini’s logical hallucinations fall sharply from 248 to 143 (−105), which is the clearest manifestation of this mechanism. When the model is forced to derive rather than guess, many problems that cannot be solved in a single step become answerable.

Visual grounding. CoT prompts require the model to explicitly cite visual evidence during reasoning, reducing its tendency to take linguistic-prior shortcuts. This aligns with the visual-grounding motivation in Zheng et al. (2024). Gemini’s faithfulness hallucinations fall from 278 to 212 because the reasoning constraint pushes the model to inspect the chart before describing it.

Exposure effect. Visual perception errors that remain hidden in Direct mode are explicitly described in the CoT reasoning chain and captured by the MLLM Judge. Even when the final answer is correct, visual misreadings in intermediate steps lower the score. This mechanism explains why stronger-baseline models see increased HR despite high final-answer accuracy.

Commitment error. CoT forces the model to make explicit visual or mathematical assertions early in the reasoning process. Once an intermediate step goes wrong, the model continues reasoning around the incorrect premise, producing “informative but wrong” responses. Once the model produces a hallucination early in reasoning, it elaborates coherently around the erroneous premise, causing the error to propagate and amplify along the chain.

These four mechanisms determine CoT’s two-sided effect. When baseline reasoning ability is weak, structured decomposition and visual grounding outweigh the risks of exposure and commitment. When the baseline is already strong and the single-step answer error rate is low, the remaining room for improvement is small, but every additional reasoning step introduces new opportunities for exposure effects and commitment errors.

5.3 Failure Case Analysis

This section presents two typical failure cases, corresponding to commitment error and exposure effect respectively.

5.3.1 Case 1: CoT Commitment Error Leading to Logical Hallucination (GPT-5.4-mini, PID=30)

Table 5.2 presents the sample information, and Figure 5.3 shows the accompanying image. The geometry problem asks for the angle in an intersecting-chords diagram. The correct answer is 70° .

In Direct mode, the model directly outputs the correct option and receives score 5. In CoT mode, the model attempts step-by-step reasoning, but makes an incorrect commitment in the first step: it misapplies the triangle angle-sum formula. After finding that the computed answer

Table 5.2: MathVista Case 1: sample information.

Field	Content
Question type	Geometry reasoning (multiple choice)
Question	Two chords AB , CD intersect at E ; $\angle D = 35^\circ$. Find $\angle C$.
Ground Truth	70° (B)




Figure 5.3: MathVista Case 1: intersecting chords in a circle.

(40°) is not among the options, it fabricates a supplementary-angle explanation and arrives at the wrong answer 80° .

This demonstrates a snowballing hallucination pattern: once an incorrect premise is locked in, subsequent steps are built on top of it. Each step may appear plausible in isolation, but the entire reasoning chain rests on the initial false assumption. When the model lacks the precise geometric knowledge needed to maintain a valid chain of reasoning, CoT can turn an originally correct intuition into a wrong answer.

5.3.2 Case 2: CoT Exposure of Visual Perception Error Leading to Faithfulness Hallucination (Qwen3-VL-235B, PID=15)

Table 5.4 presents the sample information, and Figure 5.4 shows the accompanying image. The question asks which organism would be most affected if algae were eliminated from a food web. The correct answer is Common water flea (option B).

In Direct mode the model outputs the correct answer and receives score 5. In CoT mode the model provides a detailed food web analysis but incorrectly claims that “Algae has arrows pointing to Tilapia.” In fact, the diagram contains no such arrow. This is a typical exposure effect: in Direct mode the model may reach the correct answer through some shortcut (linguistic prior or partial visual perception) while its visual perception is not fully accurate. CoT forces the model to explicitly describe the arrow relationships in the image, exposing the visual misreading to the MLLM Judge and causing the score to drop from 5 to 2.

This case reveals an important phenomenon in CoT evaluation: a correct answer does not imply correct reasoning.

Table 5.3: MathVista Case 1: responses under two modes.

Mode	Response	Score	Type
Direct	B (70°)	5	none
CoT	Flawed reasoning, arrives at 80° (C)	1	logical

Table 5.4: MathVista Case 2: sample information.

Field	Content
Question type	Biology food web reasoning (multiple choice)
Question	Which organism will be most affected if algae is eliminated?
Ground Truth	Common water flea (B)




Figure 5.4: MathVista Case 2: food web diagram.

6 Extension: VQA-RAD Medical VQA

Medical visual question answering is a representative application of MLLMs in high-stakes domains. Compared with general VQA and mathematical reasoning, medical VQA imposes higher demands on both visual perception and domain knowledge, requiring models to correctly identify anatomical structures and pathological findings in X-ray, CT, and MRI images, understand the clinical semantics of questions, and provide diagnosis-related answers grounded in visual evidence. This section evaluates the four models on VQA-RAD (Lau et al., 2018), a standard radiology VQA benchmark.

6.1 Results

Table 6.1 reports the overall evaluation results and hallucination rates stratified by answer type.

6.2 Result Analysis

6.2.1 Absolute Dominance of Faithfulness Hallucinations

The most salient feature of VQA-RAD is that faithfulness hallucinations account for more than 96% of all hallucinations across the four models (96.6%–99.0%), while factuality and logical hallucinations are almost absent. This sharply contrasts with MathVista, where all three hallucination types are substantially represented. In medical VQA, errors almost entirely originate from visual misreading rather than knowledge retrieval or logical reasoning.

This pattern can be explained through the dual-modality prior analysis in Section 4.2.3. In the radiology setting, the visual encoder faces an image domain that is underrepresented in its training distribution: general-purpose visual encoders learn primarily from natural images such as COCO and ImageNet, and lack fine-grained sensitivity to anatomical structures and lesions in X-ray, CT, and MRI. At the same time, the language module has accumulated abundant medical knowledge from text pretraining. The models can fluently discuss DWI sequences,

Table 5.5: MathVista Case 2: responses under two modes.

Mode	Response	Score	Type
Direct	B (correct answer)	5	none
CoT	B (correct, but reasoning contains visual error)	2	faithfulness

Table 6.1: VQA-RAD evaluation results. Best values are highlighted in pink.

Overall Results				
Metric	GPT-5.4-mini	Gemini 2.5 Flash	Qwen3.5-35B	Qwen3-VL-235B
Hallucination Rate	33.26%	66.30%	39.25%	58.31%
Average Score	4.52	2.68	4.20	3.57
$n_{\text{hallucinated}}$	150	299	177	263
Faithfulness	145 (96.7%)	296 (99.0%)	171 (96.6%)	257 (97.7%)
Factuality	1 (0.7%)	2 (0.7%)	4 (2.3%)	5 (1.9%)
Logical	4 (2.7%)	1 (0.3%)	2 (1.1%)	1 (0.4%)
None	301	152	274	188
Stratified by Answer Type				
Answer Type (N)	GPT-5.4-mini	Gemini 2.5 Flash	Qwen3.5-35B	Qwen3-VL-235B
Closed (yes/no, 251)	28.69%	61.75%	31.87%	50.20%
Open (free-text, 200)	39.00%	72.00%	48.50%	68.50%

periventricular white matter, and the radiologist’s-view convention, yet they cannot reliably decide which half of the brain contains a lesion. When the visual signal is intrinsically weak, the language prior dominates the output, producing responses that are fluent in medical language but visually wrong.

6.2.2 Model Ranking Reversal and Conservative Strategy Transfer

The model ranking on VQA-RAD barely overlaps with the rankings on POPE and MathVista. Qwen3-VL-235B, which leads on POPE, regresses substantially in the medical domain (HR = 58.31%), and Qwen3.5-35B-A3B, which leads on MathVista, also drops to HR = 39.25%. By contrast, GPT-5.4-mini ranks first on VQA-RAD with the lowest hallucination rate (33.26%) and highest average score (4.52). Consistent with the conservative response tendency observed on POPE, GPT-5.4-mini’s restrained style on closed-ended screening questions unexpectedly becomes an advantage here: because it tends to avoid affirmative claims under uncertainty, it produces fewer faithfulness hallucinations caused by over-claiming.

6.3 Failure Case Analysis

6.3.1 Case 1: Conservative Refusal Leading to Missed Diagnosis (Closed Question)

Table 6.2 and Figure 6.1 present this case. The question asks whether the chest X-ray shows evidence of an aortic aneurysm, with the expected answer being “Yes.”

Table 6.2: VQA-RAD Case 1: model responses. Expected answer: Yes.

Model	Response summary	Score	Type
GPT-5.4-mini	No convincing evidence; mild aortic widening	2	faithfulness
Gemini 2.5 Flash	Limited sensitivity on chest X-ray, cannot confirm	2	faithfulness
Qwen3.5-35B	Evidence suggests descending aortic aneurysm	6	none
Qwen3-VL-235B	No clear radiographic evidence	2	faithfulness

Three of the four models (GPT-5.4-mini, Gemini 2.5 Flash, Qwen3-VL-235B) correctly describe abnormal signs such as mild widening or prominence of the mediastinal aorta, but ultimately refuse to make a positive diagnosis. In other words, they perceive the abnormality but retreat at the decision stage. Qwen3.5-35B-A3B is the only model that converts the observation




Figure 6.1: VQA-RAD Case 1: chest X-ray.

into a positive diagnosis. In clinical screening, this conservative answer is equivalent to a missed diagnosis: a false negative can prevent the patient from entering the downstream diagnostic pathway.

6.3.2 Case 2: Fluent Knowledge Overriding Faulty Vision (Open Question)

Table 6.3 and Figure 6.2 present this case. The question asks on which side of the brain a lesion is located, with the correct answer being “left.” On the axial MRI shown, a clear hyperintense periventricular lesion sits on the viewer’s right of the image, which by standard radiological convention (“radiologist’s view,” patient-right on viewer-left) corresponds to the patient’s left brain.




Figure 6.2: VQA-RAD Case 2: axial brain MRI. The hyperintense periventricular lesion sits on the viewer’s right of the image, i.e., the patient’s *left* hemisphere.

All four models answer incorrectly, but through different paths. GPT-5.4-mini visually identifies that the lesion is on the right side of the image, but applies the radiological convention in the wrong direction. The two Qwen models do the opposite: they accurately recite the radiologist’s-view rule, but mislocalize the lesion in the image. This is a typical case of P_L overriding P_V : when visual evidence is difficult to interpret, fluent clinical terminology fills the gap, producing an answer that sounds professional but is visually wrong.

Table 6.3: VQA-RAD Case 2: model responses. Expected answer: *left*.

Model	Response summary	Score	Type
GPT-5.4-mini	Sees the lesion on the image right, but reverses the radiological convention → concludes right	2	faithfulness
Gemini 2.5 Flash	Misreads the brain MRI as an abdominal CT; describes gallstones in the gallbladder	0	faithfulness
Qwen3.5-35B	Recites the convention correctly, but locates the lesion on the image left → concludes right	2	faithfulness
Qwen3-VL-235B	Recites the convention and adds DWI/ischemia details, but locates the lesion in the right frontal lobe	2	faithfulness

7 Error Analysis

7.1 Human Alignment Verification (Human-as-Judge)

7.1.1 Experimental Setup

This experiment examines the agreement between the GPT-5.5 Judge’s hallucination judgments and those of a human annotator, following the human alignment methodology of [Zhou et al. \(2026\)](#). Evaluation metrics include Accuracy, F1, Cohen’s κ (chance-corrected agreement), and Type Agreement (the proportion of matching hallucination type labels among samples that both human and judge classify as hallucinated). κ values are interpreted on the conventional scale: <0.20 slight, 0.21–0.40 fair, 0.41–0.60 moderate, 0.61–0.80 substantial, 0.81–1.00 almost perfect.

Forty samples are drawn from MathVista testmini CoT results via stratified sampling (10 per model, approximately 5 hallucinated and 5 non-hallucinated each, seed 42), prioritizing coverage of all three hallucination types within the hallucinated group. The annotation interface displays each image, question, model response, and reference answer. The annotator judges whether hallucination is present and, if so, labels the type (faithfulness, factuality, or logical). A single-blind design (model name and judge verdict hidden) is used to avoid annotation bias. Figure 7.1 shows the annotation interface.

7.1.2 Overall Results

Substantial agreement between automatic Judge and human annotator. On the 40 valid human-annotated samples, the GPT-5.5 Judge reaches 85.0% accuracy (95% CI: [0.709, 0.929]), $F1 = 0.824$, and Cohen’s $\kappa = 0.700$ (95% CI: [0.481, 0.898]). Under the conventional interpretation, $\kappa = 0.700$ falls into the substantial-agreement range, indicating that the automatic Judge tracks human judgments closely. Table 7.1 reports the overall and per-model breakdowns.

Table 7.1: Human alignment results. Per-model intervals are wide because each model is evaluated on only 10 samples.

Subset	N	Accuracy [95% CI]	Precision [95% CI]	Recall [95% CI]	F1	Cohen’s κ [95% CI]
Overall	40	0.850 [0.709, 0.929]	0.700 [0.481, 0.855]	1.000 [0.785, 1.000]	0.824	0.700 [0.481, 0.898]
Qwen3-VL-235B	10	0.800 [0.490, 0.943]	0.600 [0.231, 0.882]	1.000 [0.439, 1.000]	0.750	0.600 [0.000, 1.000]
Qwen3.5-35B	10	1.000 [0.722, 1.000]	1.000 [0.566, 1.000]	1.000 [0.566, 1.000]	1.000	1.000 [1.000, 1.000]
Gemini 2.5 Flash	10	0.600 [0.313, 0.832]	0.200 [0.036, 0.624]	1.000 [0.207, 1.000]	0.333	0.200 [0.000, 0.737]
GPT-5.4-mini	10	1.000 [0.722, 1.000]	1.000 [0.566, 1.000]	1.000 [0.566, 1.000]	1.000	1.000 [1.000, 1.000]

Errors concentrate on over-judgment, not missed detection. The confusion matrix (Table 7.2) contains 20 true negatives, 14 true positives, 6 false positives, and 0 false negatives.




Figure 7.1: Human annotation interface used in the alignment experiment. Each screen shows the image, question, model response, and reference answer, with model name and judge verdict hidden to avoid bias.


The Judge therefore achieves a perfect recall of 1.000 (95% CI: [0.785, 1.000]): every response that the human annotator labels as hallucinated is also flagged by the Judge. The dominant error mode is over-judgment rather than missed detection.

Table 7.2: Confusion matrix: human annotation (rows) vs. GPT-5.5 Judge verdict (columns).

Human \ Judge	No halluc.	Halluc.	Total
No halluc.	20	6	26
Halluc.	0	14	14
Total	20	20	40

All six disagreements fall into the “human-no, Judge-yes” quadrant. They concentrate on items that ask the model to infer information not directly visible in the image, such as exact age differences or birth decades. In these samples the model typically declines to estimate from appearance alone and outputs a conservative answer of 0 or “cannot determine”. The human annotator treats such a refusal as no additional hallucination, because the model explicitly avoids fabricating unobservable content, whereas the Judge tends to mark it as a logical or faithfulness hallucination solely on the basis of the numerical mismatch with the reference answer. Figure 7.2 shows three representative cases.

Hallucination type labels also align well. Type agreement reaches 0.9286: among the 14 samples that both human and Judge classify as hallucinated, 13 receive identical type labels, with only one logical-vs-faithfulness mismatch. This confirms that the three-category taxonomy proposed in this report is operationally well-defined: the Judge not only achieves high recall in detecting hallucination but also assigns type labels that are consistent with the human annotation in most cases. Per-model confidence intervals nonetheless remain wide, because this study relies on a single annotator and only 40 samples (10 per evaluated model); larger annotation sets and multiple annotators would be needed for finer-grained per-model reliability claims.



(a) Age gap between the leftmost and rightmost persons (reference 5). The model refuses to infer ages from appearance and answers 0.

(b) Age gap between the two depicted persons (reference 27). The model explicitly invokes the “visible-information-only” rule and answers 0.

(c) Number of people born after the end of WWII (reference 2). The model cannot determine birth years from the image and answers 0.

Figure 7.2: Three representative over-judgment cases in the human alignment study (annotation interface screenshots). In every case the model declines to infer unobservable information and outputs a conservative 0; the human annotator labels them as no hallucination, while the Judge flags them based on the numerical mismatch with the reference answer.

7.2 Judge Model Consistency Check

Setup. This experiment examines whether evaluation conclusions depend on the choice of judge model, by comparing GPT-5.5 and Claude Opus 4.7 on the same sample set. One hundred samples are randomly drawn from MathVista testmini (seed 42), yielding 100 responses per evaluated model (400 total). Both judges use identical system prompts (MMHal-Bench 0–6 rubric, five few-shot examples, hallucination type classification).

Table 7.3 reports the results; the “Model” column refers to the model whose responses are being judged, not to the judge.

Table 7.3: Judge consistency: GPT-5.5 vs. Claude Opus 4.7. The “Model” column denotes the response model being judged, not the judge model.

Model	N	GPT HR	Claude HR	Δ HR	GPT Avg	Claude Avg	κ	Spearman ρ
GPT-5.4-mini	100	45.0%	44.0%	−1.0 pp	3.27	3.01	0.980	0.805
Gemini 2.5 Flash	100	37.0%	27.0%	−10.0 pp	4.30	4.40	0.591	0.658
Qwen3.5-35B	100	14.0%	18.0%	+4.0 pp	5.01	4.61	0.777	0.722
Qwen3-VL-235B	100	15.0%	22.0%	+7.0 pp	4.91	4.38	0.638	0.669
Overall	400	n/a	n/a	n/a	n/a	n/a	0.776	0.740

Note: Overall κ and ρ are pooled across the 400 samples, not arithmetic means of the per-model values. Overall Pearson $r = 0.762$ ($p < 0.001$).

Substantial agreement with preserved ranking. The two judges show substantial overall agreement ($\kappa = 0.776$, Spearman $\rho = 0.740$), and the relative model ranking is preserved under both: in each judge, both Qwen models rank best, GPT-5.4-mini sits in the middle, and Gemini 2.5 Flash is the weakest. The largest per-model divergence falls on Gemini ($\kappa = 0.591$): Claude rates Gemini’s HR 10 pp lower than GPT-5.5 does, indicating that Gemini’s verbose responses sit closer to the disagreement boundary between the two judges’ calibration of “informative but wrong.” GPT-5.4-mini is by contrast the most stable ($\kappa = 0.98$), likely because its often single-token Direct responses leave little room for judge interpretation. These conclusions are based on two frontier closed-source judges and should be re-examined when an open-source or domain-specific judge becomes available.

7.3 Threshold Sensitivity Analysis

Setup. This analysis examines how sensitive the Hallucination Rate is to the threshold $T = 3$. For each threshold T from 0 to 6, $HR(T)$ is computed as $N_{\text{score} < T} / N_{\text{total}}$.

Table 7.4 reports the results for three representative thresholds.

Table 7.4: Threshold sensitivity (MathVista Direct).

Model	HR($T=2$)	HR($T=3$)	HR($T=4$)	$\Delta(3 \rightarrow 2)$	$\Delta(4 \rightarrow 3)$
GPT-5.4-mini	18.1%	45.1%	45.3%	−27.0 pp	+0.2 pp
Gemini 2.5 Flash	1.2%	34.3%	34.6%	−33.1 pp	+0.3 pp
Qwen3.5-35B	0.8%	13.3%	13.3%	−12.5 pp	+0.0 pp
Qwen3-VL-235B	2.1%	16.3%	16.4%	−14.2 pp	+0.1 pp




Figure 7.3: Hallucination rate as a function of threshold T (MathVista Direct). The vertical dashed line marks the official threshold $T = 3$.

Bimodal scores; ranking robust to threshold. $HR(T=3)$ and $HR(T=4)$ are nearly identical (difference < 0.3 pp), which implies that very few samples receive score 3 and the judge’s scores follow a clear bimodal distribution. The largest jump occurs between $T=2$ and $T=3$ (12.5–33.1 pp), indicating a concentration of samples at score 2. The model ranking is fully preserved across all three thresholds, so the core conclusions are insensitive to the threshold choice.

7.4 Length Bias Analysis

Setup. This analysis examines whether the MLLM Judge exhibits a systematic length bias. Spearman correlation coefficients between response length (word count) and judge score are computed for all MathVista samples in both Direct and CoT modes. Each row also reports the median length per mode and the mean score across length quartiles (Q1 = shortest, Q4 = longest).

Table 7.5 reports the results.

No uniform length bias; CoT shifts the sign. The judge does not show a uniform “longer $\rightarrow$ higher score” bias: the sign of the correlation depends on the quality pattern of each model’s responses. The correlation is positive when longer responses carry more valid reasoning (both Qwen models in Direct mode), and negative when longer responses introduce more errors (Gemini, and Qwen3-VL under CoT). The Qwen3.5 Direct correlation of +0.530

Table 7.5: Spearman correlation between response length and judge score (MathVista). Q1→Q4 columns give the mean score across length quartiles (Q1 = shortest 25%, Q4 = longest 25%). **: $p < 0.01$; ***: $p < 0.001$.

Model	Direct ρ	CoT ρ	Direct median	CoT median	Direct Q1→Q4	CoT Q1→Q4
GPT-5.4-mini	+0.086**	+0.015	1	43	3.19→3.46	4.17→3.95
Gemini 2.5 Flash	−0.142***	−0.218***	242	228	4.61→3.80	5.30→4.33
Qwen3.5-35B	+0.530***	+0.006	42	122	4.56→5.33	5.32→5.26
Qwen3-VL-235B	+0.394***	−0.241***	40	103	4.38→4.92	5.32→4.32

is the strongest in the panel, but the corresponding Q1→Q4 progression (4.56 → 5.33) shows that the longer-quartile responses do score better in absolute terms, indicating “longer because more correct” rather than “higher score because longer.”

Switching to CoT generally eliminates or reverses the positive correlation: Qwen3.5 drops from +0.530 to nearly zero, and Qwen3-VL flips from +0.394 to −0.241. This pattern is consistent with the exposure effect identified in Section 5: under CoT, longer chains contain more visual assertions and therefore more opportunities to be penalized, so the judge becomes more sensitive to noise in longer reasoning chains.

7.5 Cross-Model Hallucination Consistency

Setup. This analysis examines whether hallucinations overlap across models on the same questions. For each sample, the number of models that produce hallucinations (0–4) is counted.

Table 7.6 reports the distribution.

Table 7.6: Cross-model hallucination overlap. Hard Ratio = proportion of samples where ≥ 3 out of 4 models hallucinate.

Dataset	N	0/4	1/4	2/4	3/4	4/4	Hard Ratio
MathVista Direct	1000	387	321	175	49	68	11.7%
MathVista CoT	1000	487	232	132	77	72	14.9%
VQA-RAD	451	59	118	122	81	71	33.7%

VQA-RAD difficulty is largely intrinsic. On VQA-RAD, only 13.1% of samples (59/451) are answered correctly by all four models, and 33.7% are answered incorrectly by three or more. This concentration at the “hard” end of the spectrum indicates a substantial pool of items whose difficulty is intrinsic to the data rather than specific to any individual model. The finding sharpens the interpretation of Section 6: the high VQA-RAD hallucination rates are not driven by an idiosyncratic weakness of one or two models but by a domain-level perception gap shared across general-purpose MLLMs, consistent with the “knowledge $\neq$ perception” diagnosis of Case 2 (Section 6.3.2).

MathVista and the polarization signature of CoT. MathVista Direct shows a more balanced distribution (38.7% fully correct, 6.8% fully incorrect). Under CoT the Hard Ratio rises from 11.7% to 14.9%, a modest but consistent shift in line with the CoT polarization effect already documented in Section 5: CoT pushes responses toward the extremes of either fully-correct or fully-wrong, so the distribution at both tails grows while intermediate cells shrink.




Figure 7.4: Radar chart comparing model performance across six dimensions (accuracy = 1 – HR). Points closer to the outer ring indicate better performance.

8 Limitations and Future Work

8.1 Limitations

Benchmark coverage is narrow. The three datasets (POPE, MathVista, VQA-RAD) probe three distinct hallucination regimes (object existence, mathematical reasoning, and medical visual grounding) but do not span the full landscape of publicly available hallucination benchmarks. HallusionBench (Guan et al., 2024) (visual illusions and language-context manipulations), AMBER (Wang et al., 2023) (a multi-dimensional LLM-free benchmark covering existence, attribute, and relation hallucinations), and broader capability suites such as MM-Vet (Yu et al., 2024) and MMBench (Liu et al., 2024b) are not included. Whether the patterns identified here (particularly the CoT polarization effect on MathVista and the knowledge–perception gap on VQA-RAD) generalize to these probes remains untested.

Model panel is restricted. Four models from three families (OpenAI, Google, Alibaba) are evaluated. Smaller open-weight models (e.g., the LLaVA, MiniGPT, and InstructBLIP families) and parameter-scale ablations within a single family are excluded, so the report cannot disentangle effects driven by model scale from those driven by training methodology or visual-encoder choice.

Human alignment is sample-limited. The 40-sample annotation pool yields an overall Cohen’s $\kappa = 0.70$ with a 95% bootstrap CI of $[0.48, 0.90]$, which is wide enough to characterize agreement as substantial in aggregate but insufficient to support a statistically grounded per-model judge-reliability ranking. A scaled-up annotation effort (approximately $N \geq 80$ per model) would be required to make such quantitative cross-model claims.

8.2 Future Work

Language-prior override of visual content. Recently, in a casual conversation with GPT-5.5, I noticed an interesting phenomenon: when I asked “Which orange ball is bigger?” on the

pair of Ebbinghaus-style images shown in Figure 8.1, the model gave the same answer on both inputs, even though the second image is not a real illusion at all. What I find striking is that this is essentially the same mechanism identified in the POPE Popular split (Section 4.3.2), where models fall back on co-occurrence statistics once a familiar context cue is present; the Ebbinghaus probe just pushes it to an extreme, in which the memorized association “Ebbinghaus image $\Rightarrow$ equal sizes” appears to override the image signal entirely. This is consistent with the reasoning-time visual disengagement reported by Liu et al. (2025) and the broader language-prior failure modes surveyed by Bai et al. (2025b).




Figure 8.1: Conversation screenshots from a manual probe of GPT-5.5 on two Ebbinghaus-style images. The model gives identical answers on both inputs despite the obvious size difference in the right image.

A LAST-ViT-inspired remediation. I happened to have read a recent paper that addresses a structurally similar problem on the pure-vision side. Shi et al. (2026) observe that in standard Vision Transformers, the model’s summary representation (the CLS token, which collects information from all image patches and is then used to make the final prediction) tends to draw most of its content from background regions of an image rather than from the actual subject of the photo. The reason is fairly mundane: background patches simply occupy more area than foreground patches, and the standard training objective only checks whether the final image-level label is correct, without constraining which patches the model is actually paying attention to. Their proposed fix, named LAST-ViT, is a lightweight plug-in module that re-weights how patch information is aggregated into the CLS token, biasing it back toward the foreground without requiring any retraining of the backbone, and the paper reports consistent gains across twelve vision benchmarks. The qualitative effect of the fix is illustrated in Figure 8.2: a vanilla Transformer scatters attention across background patches, whereas the LAST-ViT variant concentrates attention on the actual subject of the photo.

Reading that paper alongside my own Ebbinghaus observation, I started wondering whether the same kind of shortcut is at work in MLLMs as well, just along a different axis: instead of background patches drowning out foreground patches in a single image, the language prior drowns out the image signal across an entire conversation turn. The dual-prior framework of CausalMM (Section 4.2.3) provides a causal grounding for this conjecture: P_V and P_L act through independent confounding pathways, and CausalMM showed that targeted intervention on either pathway can partially offset the other’s confounding effect. Following the LAST-ViT approach, a decoding-time re-weighting strategy that pulls attention back toward visual tokens when the two priors conflict may improve model performance. A concrete next step would be to study whether such re-weighting at decoding time can reduce this failure mode without




Figure 8.2: Attention and feature visualizations comparing a ConvNet baseline, a standard Vision Transformer, and a Transformer equipped with the LAST-ViT (LazyStrike) module. The Transformer baseline disperses attention across the background, while the LAST-ViT variant concentrates attention on the foreground subject. Figure reproduced from Shi et al. (2026), Fig. 1.

degrading language capability.

References

- Anthropic. Claude opus 4.7 model card. <https://www.anthropic.com/claude>, 2026. Accessed: 2026-05-20.
- Shuai Bai, Yuxuan Cai, Ruizhe Chen, et al. Qwen3-vl technical report. arXiv preprint arXiv:2511.21631, 2025a.
- Zechen Bai, Pichao Wang, Tianjun Xiao, Tong He, Zongbo Han, Zheng Zhang, and Mike Zheng Shou. Hallucination of Multimodal Large Language Models: A Survey. arXiv preprint arXiv:2404.18930, 2025b.
- Zhiyuan Chen, Yuecong Min, Jie Zhang, Bei Yan, Jiahao Wang, Xiaozhen Wang, and Shiguang Shan. A Survey of Multimodal Hallucination Evaluation and Detection. *International Journal of Computer Vision*, 2026. Online First.
- Google DeepMind. Gemini api models. <https://ai.google.dev/gemini-api/docs/models>, 2026. Accessed: 2026-05-20.
- Tianrui Guan, Fuxiao Liu, Xiyang Wu, Ruiqi Xian, Zongxia Li, Xiaoyu Liu, Xijun Wang, Lichang Chen, Furong Huang, Yaser Yacoob, Dinesh Manocha, and Tianyi Zhou. HallusionBench: An Advanced Diagnostic Suite for Entangled Language Hallucination and Visual Illusion in Large Vision-Language Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Sai Srinivas Kancheti, Aditya Sanjiv Kanade, Vineeth N. Balasubramanian, and Tanuja Ganu. Chain-of-Thought Degrades Visual Spatial Reasoning Capabilities of Multimodal LLMs. In *Proceedings of the Conference on Language Modeling (COLM) 2025*, 2025.
- Jason J. Lau, Soumya Gayen, Asma Ben Abacha, and Dina Demner-Fushman. A dataset of clinically generated visual questions and answers about radiology images. *Scientific Data*, 5: 180251, 2018.

- Yifan Li, Yifan Du, Kun Zhou, Jinpeng Wang, Wayne Xin Zhao, and Ji-Rong Wen. Evaluating Object Hallucination in Large Vision-Language Models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 292–305, 2023.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C. Lawrence Zitnick. Microsoft COCO: Common Objects in Context. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2014.
- Chengzhi Liu, Zhongxing Xu, Qingyue Wei, Juncheng Wu, James Zou, Xin Eric Wang, Yuyin Zhou, and Sheng Liu. More thinking, less seeing? Assessing Amplified Hallucination in Multimodal Reasoning Models. In *Advances in Neural Information Processing Systems (NeurIPS) 2025*, 2025.
- Hanchao Liu, Wenyuan Xue, Yifei Chen, Dapeng Chen, Xiutian Zhao, Ke Wang, Liping Hou, Rongjun Li, and Wei Peng. A Survey on Hallucination in Large Vision-Language Models. arXiv preprint arXiv:2402.00253, 2024a.
- Yuan Liu, Haodong Duan, Yuanhan Zhang, Bo Li, Songyang Zhang, Wangbo Zhao, Yike Yuan, Jiaqi Wang, Conghui He, Ziwei Liu, Kai Chen, and Dahua Lin. MMBench: Is Your Multimodal Model an All-around Player? In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2024b.
- Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi, Hao Cheng, Kai-Wei Chang, Michel Galley, and Jianfeng Gao. MathVista: Evaluating Mathematical Reasoning of Foundation Models in Visual Contexts. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024.
- OpenAI. Gpt-5.5 model card. <https://platform.openai.com/docs/models>, 2026a. Accessed: 2026-05-20.
- OpenAI. Openai models overview. <https://platform.openai.com/docs/models>, 2026b. Accessed: 2026-05-20.
- Qwen Team. Qwen3.5-35b-a3b: Multimodal mixture-of-experts model. <https://huggingface.co/Qwen/Qwen3.5-35B-A3B>, 2025a. Accessed: 2026-05-20.
- Qwen Team. Qwen3-vl-235b-a22b-instruct: Vision-language moe model. <https://huggingface.co/Qwen/Qwen3-VL-235B-A22B-Instruct>, 2025b. Accessed: 2026-05-20.
- Anna Rohrbach, Lisa Anne Hendricks, Kaylee Burns, Trevor Darrell, and Kate Saenko. Object Hallucination in Image Captioning. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2018.
- Cheng Shi, Yizhou Yu, and Sibe Yang. Vision Transformers Need More Than Registers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026.
- Zhiqing Sun, Sheng Shen, Shengcao Cao, Haotian Liu, Chunyuan Li, Yikang Shen, Chuang Gan, Liang-Yan Gui, Yu-Xiong Wang, Yiming Yang, Kurt Keutzer, and Trevor Darrell. Aligning Large Multimodal Models with Factually Augmented RLHF. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 13088–13110, 2024.
- Junyang Wang, Yuhang Wang, Guohai Xu, Jing Zhang, Yukai Gu, Haitao Jia, Ming Yan, Ji Zhang, and Jitao Sang. AMBER: An LLM-free Multi-dimensional Benchmark for MLLMs Hallucination Evaluation. arXiv preprint arXiv:2311.07397, 2023.

- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Ed Chi, Quoc Le, and Denny Zhou. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS) 2022*, 2022.
- Jinge Wu, Yunsoo Kim, and Honghan Wu. Hallucination Benchmark in Medical Visual Question Answering. In *The Second Tiny Papers Track at ICLR 2024*, 2024.
- Weihao Yu, Zhengyuan Yang, Linjie Li, Jianfeng Wang, Kevin Lin, Zicheng Liu, Xinchao Wang, and Lijuan Wang. MM-Vet: Evaluating Large Multimodal Models for Integrated Capabilities. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.
- Haojie Zheng, Tianyang Xu, Hanchi Sun, Shu Pu, Ruoxi Chen, and Lichao Sun. Thinking before looking: Improving Multimodal LLM Reasoning via Mitigating Visual Hallucination. arXiv preprint arXiv:2411.12591, 2024.
- Weihong Zhong, Xiaocheng Feng, Liang Zhao, Qiming Li, Lei Huang, Yuxuan Gu, Weitao Ma, Yuan Xu, and Bing Qin. Investigating and mitigating the multimodal hallucination snowballing in large vision-language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 11991–12011, 2024.
- Guanyu Zhou, Yibo Yan, Xin Zou, Kun Wang, Aiwei Liu, and Xuming Hu. Mitigating modality prior-induced hallucinations in multimodal large language models via deciphering attention causality. In *Proceedings of the 13th International Conference on Learning Representations (ICLR)*, 2025.
- Jiawei Zhou, Chi Zhang, Xiang Feng, Qiming Zhang, Haibo Qiu, Lihuo He, Dengpan Ye, Xinbo Gao, and Jing Zhang. Omni-I2C: A Holistic Benchmark for High-Fidelity Image-to-Code Generation. arXiv preprint arXiv:2603.17508, 2026.

A Dataset Details and Prompt Templates

A.1 Dataset Details

A.1.1 POPE (Polling-based Object Probing Evaluation)

Overview. POPE (Li et al., 2023) converts object hallucination evaluation into a binary existence probing task built on top of the MS COCO val2014 split (Lin et al., 2014). The paradigm was originally proposed by Rohrbach et al. (2018) under the CHAIR metric for image captioning hallucination, and later systematized into a stable polling protocol by Li et al. (2023).

Construction. Given a COCO image with annotated objects, POPE samples six probing questions per image: three positive ones (asking about a ground-truth-present object) and three negative ones (asking about a ground-truth-absent object). Each question follows the fixed template “Is there a/an <object> in the image?” and expects a Yes/No answer. The three splits differ only in how the negative objects are sampled:

- *Random:* negative objects are drawn uniformly from the COCO label set, excluding objects present in the image.
- *Popular:* negative objects are drawn from the top-K most frequent COCO categories, again excluding the present ones, creating high-frequency co-occurrence distractors.
- *Adversarial:* for each image, the three categories that co-occur most frequently with the present objects in the COCO training statistics are chosen as negatives, yielding semantically similar distractors.

The three splits induce a monotonic difficulty gradient, empirically confirmed in Section 4.2.2.

Scale and balance. Each split contains 3,000 probing questions (500 images $\times$ 6 questions), with a balanced 1,500 Yes / 1,500 No ground-truth distribution. Official evaluation reports Accuracy, Precision, Recall, F1, and Yes Ratio.

Subset used in this experiment. For each of the three splits, 1,000 questions are randomly subsampled with seed 42, yielding 3,000 evaluation samples per model and 12,000 in total across the four evaluated models. The 1,000-per-split budget balances API-cost and statistical power while preserving the original 1:1 Yes/No balance of the parent splits.

A.1.2 MathVista

Overview. MathVista (Lu et al., 2024) is a comprehensive benchmark for evaluating the mathematical reasoning of foundation models in visual contexts. The task requires simultaneous visual perception, mathematical knowledge, and multi-step reasoning, making it well suited to trigger and analyze all three hallucination types (faithfulness, factuality, logical).

Construction. The benchmark integrates 31 source datasets: 28 existing multimodal math or VQA datasets and three datasets newly created by the authors (IQTest, FunctionQA, PaperQA). Each sample is annotated with two orthogonal taxonomies:

- *Five task types:* Figure Question Answering (FQA), Geometry Problem Solving (GPS), Math Word Problem (MWP), Textbook Question Answering (TQA), and Visual Question Answering (VQA).
- *Seven mathematical-reasoning skills:* algebraic, arithmetic, geometric, logical, numeric commonsense, scientific, and statistical reasoning.

Questions appear in two answer formats: multiple choice and free-form (integer, float, or short text).

Scale. The benchmark contains 6,141 examples in total, partitioned into a 1,000-example TESTMINI subset (with public ground truth, used for model development) and a 5,141-example TEST subset (with private ground truth, used for leaderboard ranking). Official evaluation reports accuracy as the primary metric.

Subset used in this experiment. The complete TESTMINI subset (1,000 samples) is used, under both Direct and CoT prompting modes, yielding 2,000 evaluation samples per model and 8,000 in total across the four evaluated models. Both modes share the same sample set, which allows direct paired comparison in Section 5.

A.1.3 VQA-RAD

Overview. VQA-RAD (Lau et al., 2018) is the first manually constructed medical VQA dataset in which clinicians generated questions and reference answers directly from radiology images sourced from the MedPix open teaching collection. It serves as a high-stakes domain testbed because both visual perception (anatomy, pathology) and clinical knowledge are required.

Construction. Radiology trainees and faculty were asked to write naturally occurring questions about a balanced subset of radiology images. The annotation produced 3,515 question–answer pairs over 315 images, with the image collection deliberately balanced across modalities (approximately 104 head axial CT/MRI, 107 chest X-rays, and 104 abdominal axial CT slices). Each question is labeled with one of eleven categories: modality, plane, organ system, abnormality, object/condition presence, positional reasoning, color, size, attribute (other), counting, and other. Answers are split into free-form (43.1%, 1,515 pairs) and closed-ended (56.9%, 2,000 pairs); within the closed-ended portion, approximately 92% are Yes/No questions.

Scale and the standard test split. The official test split consists of 451 QA pairs: 300 randomly chosen free-form questions plus 151 paraphrased questions derived from those free-form items. Official evaluation typically reports exact-match accuracy or BLEU for free-form answers.

Subset used in this experiment. The complete official 451-item test split is used, accessed via the Hugging Face mirror of the dataset. The closed/open ratio of this test split (251 closed-ended, 200 open-ended in the HF version) allows the stratified HR reporting in Section 6. No additional sampling or filtering is applied.

A.2 Response Generation Prompts

A.2.1 POPE

The POPE prompt concatenates the dataset question with a fixed format instruction that constrains the response to a single token (Yes or No). This matches the post-processing assumption of Algorithm 1. Figure A.1 shows a concrete example.



Prompt sent to the model:

Is there a chair in the image?
Please answer YES or NO without an explanation.

Figure A.1: POPE prompt example (COCO image with a chair-related probing question).

A.2.2 MathVista

MathVista prompts are dynamically constructed from sample metadata. The format instruction varies by question type: “Answer with the option letter only.” for multiple-choice questions, “Answer with an integer.” for integer answers, “Answer with a number rounded to k decimal place(s).” for float answers, and “Keep your answer concise.” as the default. If the sample carries a non-empty `unit` field, the suffix “Include the unit ‘{unit}’.” is appended to the integer and float instructions. If choice options are present, they are appended on a separate line.

In Direct mode, the prompt structure is:

```
Question: {question}

{format_instruction}

Choices: {choices}
```

In CoT mode, a step-by-step reasoning preamble is prepended:


```
Please generate a step-by-step answer, base your reasoning strictly on what is visible
in the image, and avoid speculating about details that are unclear or not present
. End with the final answer on a separate line starting with 'Final answer:'.

Question: {question}

{format_instruction}

Choices: {choices}
```

Figure A.2 shows a concrete CoT-mode example using a geometry problem from MathVista testmini (pid=5, question_type=multi_choice).



CoT prompt sent to the model:

```
Please generate a step-by-step answer, base your
reasoning strictly on what is visible in the image,
and avoid speculating about details that are unclear
or not present. End with the final answer on a
separate line starting with 'Final answer:'.

Question: Find  $m\angle H$ 

Answer with the option letter only.
Choices: 97, 102, 107, 122
```

Figure A.2: MathVista CoT prompt example (geometry problem, pid=5).


A.2.3 VQA-RAD

The VQA-RAD prompt prepends an expert-role preamble (“You are an expert radiologist.”), following the role-setting strategy of Wu et al. (2024), and appends an instruction to ground the answer in the medical image. Figure A.3 shows a concrete example.

A.3 MLLM Judge Prompts

A.3.1 System Prompt

The judge system prompt consists of three parts: a rating scale definition, five few-shot examples with an important reminder, and a hallucination type classification with output format. The



Prompt sent to the model:

You are an expert radiologist.

Is there evidence of an aortic aneurysm?
Please provide a detailed answer based on what you observe in the medical image.

Figure A.3: VQA-RAD prompt example (chest X-ray).

complete text reproduced below matches the implementation in `evaluation/judge.py`; non-ASCII em dashes and arrows in the source are transcribed as -- and ->.

```
You are an impartial evaluator for multimodal language models. Your task is to rate a
model's response to a visual question on a 0-6 scale, jointly considering
informativeness and hallucination.

## Rating Scale (0-6)
- **6** : Very informative, with good analysis/reasoning, NO hallucination
- **5** : Very informative, NO hallucination
- **4** : Somewhat informative, NO hallucination
- **3** : Not informative (e.g., "I don't know"), but NO hallucination
- **2** : Very informative, WITH hallucination
- **1** : Somewhat informative, WITH hallucination
- **0** : Not informative, WITH hallucination

## Hallucination Definition
Hallucination refers to situations where the response includes information not present
or implied in the image, or makes claims inconsistent with the visual content. A
response that honestly admits inability to answer (e.g., "I cannot determine") has
NO hallucination (score 3).

## Few-Shot Examples

**Example 1 -- Hallucination (Score 2):**
Image shows two cats. Question: "How many cats are in the image?" Model response: "
There is one cat in the image." -> The response is informative but miscounts a
visible object. Score: 2 (Very informative, WITH hallucination).

**Example 2 -- No Hallucination (Score 5):**
Image shows a panda sitting on a tree branch with green leaves. Question: "What is the
panda doing?" Model response: "The panda is sitting on a tree branch, surrounded
by leaves." -> The response is accurate and makes no false claims about the image.
Score: 5 (Very informative, NO hallucination).

**Example 3 -- Hallucination (Score 2):**
Image shows a dog holding an umbrella, with no water in sight. Question: "Describe the
scene." Model response: "A dog holding an umbrella, positioned near the water."
-> The response fabricates "water" which is not in the image. Score: 2 (Very
informative, WITH hallucination).
```

```

**Example 4 -- No Hallucination (Score 3):**
Image shows a complex diagram or chart. Question: "Explain the economic trend." Model
response: "I cannot determine the specific trend from this image. The chart
appears to show some data over time." -> The response honestly admits uncertainty
without making false claims. Score: 3 (Not informative, NO hallucination).

**Example 5 -- No Hallucination (Score 6):**
Image shows industrial smokestacks. Question: "What environmental concerns does this
image raise?" Model response provides a detailed analysis of potential emissions,
air quality impacts, and relevant environmental regulations, all grounded in
visible evidence. -> The detailed reasoning is based on what is visible and makes
no false claims. Score: 6 (Very informative with good analysis, NO hallucination).

## Important Reminder
The standard human-generated answer may only contain factual information without
detailed analysis, and may not be fully comprehensive. A model response that
provides detailed reasoning or analysis grounded in the visual evidence should be
encouraged, not penalized. Only penalize claims that are clearly inconsistent with
or absent from the image. If the model provides extra detail that is factually
correct and consistent with the image, do NOT treat it as hallucination.

## Hallucination Types (only when score < 3)
- **faithfulness**: Visual inconsistency -- describes objects/attributes/relations not
  in the image
- **factuality**: Contradiction with established world knowledge
- **logical**: Reasoning error -- conclusion does not follow from evidence or self-
  contradiction

## Output Format
Reply in JSON only:
{
  "score": 0-6,
  "has_hallucination": true/false,
  "hallucination_type": "faithfulness" | "factuality" | "logical" | "none",
  "reason": "Brief explanation in English"
}

```

A.3.2 User Prompt

For each evaluation, the judge receives a user prompt constructed from a dataset-specific task description, the question, the model response, and the ground-truth answer. The task description varies by dataset and is selected by `_build_judge_prompt` in `evaluation/judge.py`:

```

# POPE
The question is a Yes/No object existence query.

# MathVista
The question requires mathematical visual reasoning (chart, diagram, or geometry).

# VQA-RAD
The question is about medical visual question answering based on a radiology image (X-
ray, CT, or MRI). Carefully check whether the model's response is consistent with
the visual content of the image and with established medical and anatomical
knowledge.

```


The user prompt template is:

```
{task_desc}
```

Question: {question}
 Model Response: {model_response}
 Ground Truth Answer: {ground_truth}

Please rate the model's response on the 0-6 scale.

Figure A.4 shows a concrete example using the MathVista geometry problem (pid=5), where GPT-5.4-mini produced a correct CoT response (score 6).



User prompt sent to the judge:

The question requires mathematical visual reasoning (chart, diagram, or geometry).

Question: Find $m\angle H$
 Choices: 97, 102, 107, 122
 Model Response: The figure shows a pentagon with interior angles labeled: angle E = x , angle F = $(x+20)$, angle G = $(x+5)$, angle H = $(x-5)$, angle J = $(x+10)$. For a pentagon, the sum of interior angles is 540. So $x + (x+20) + (x+5) + (x-5) + (x+10) = 540$, $5x + 30 = 540$, $x = 102$. Then $m\angle H = x - 5 = 97$.
 Final answer: A
 Ground Truth Answer: 97

Please rate the model's response on the 0-6 scale.

Figure A.4: MLLM Judge user prompt example (MathVista pid=5). The judge returns {"score": 6, "has_hallucination": false, "hallucination_type": "none", "reason": "..."}.